



# IMDb

## Internet Movie Database

Yasin Edin, Ibrahim Gökce, Norbert Jendreizik

# Content

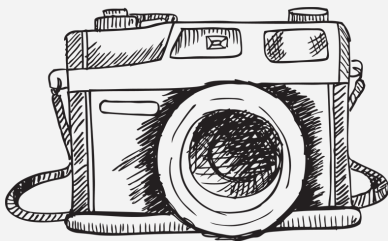
## Objectives

### Work Packages:

- Databricks
  - o Exploratory Data Analysis (EDA)
  - o Feature Engineering
- Sagemaker
  - o Model Selection
  - o Tuning

### Conclusion:

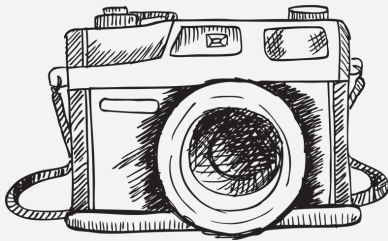
- Lessons Learned
- Future Outlook



# Objectives

# Objectives

1. **Explore** the data on a distributive computing cluster
2. **Select features** that may help you predict the ratings of movies
3. **Train** a model of your choice on your features to predict the rating of a movie
4. **Deploy** the model on AWS



# Work Packages - Databricks

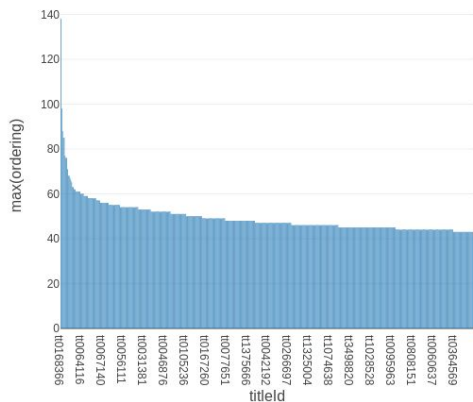
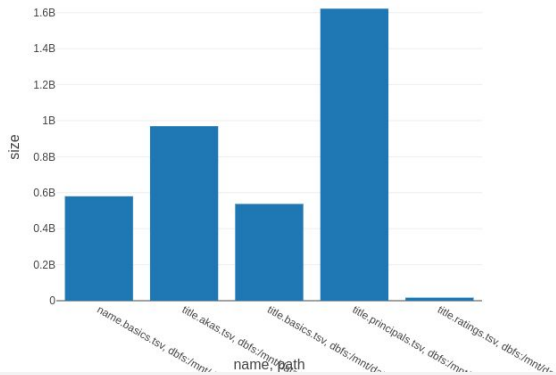
# What we had to do

## Preprocessing

- 5 tables, 5 dataframes:
  - df\_akas,
  - df\_basics,
  - df\_principals,
  - df\_ratings,
  - df\_names\*
- Cleaning files of e.g. NaN's, inconsistent values
- Encoding of e.g. region, genre etc.

## General overview

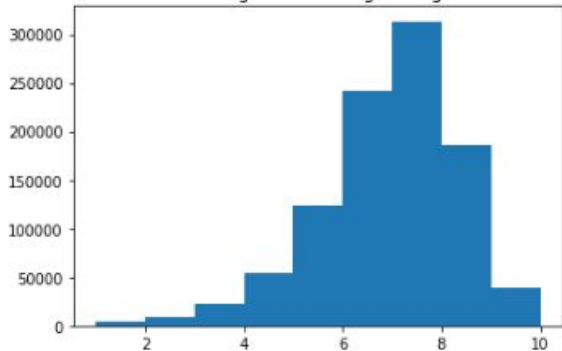
- df\_principals has the the most entries
- There are more unique titles in df\_principals than in df\_ratings



## Akas

- Contains alternative names of films. For e.g. Pokemon has one tt-number and diverse names
- Types and Attributes are more than 90% empty
- Most movies are from France, Germany, Spain, Italy, India,...

Histogram of averageRating

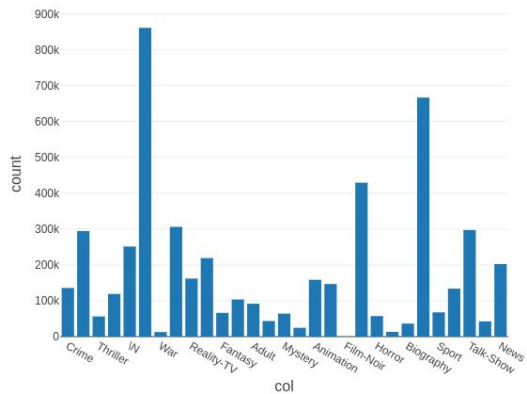


## Ratings

- The mean rating is 6.89 (maybe people are a little bit kinder than expected), with a standard deviation of roughly 1.40. We have roughly 1 million entries in total.
- 75% of the movies have 74 or fewer total votes.

## Basics

- Splitting the multiple columns into separate genres shows that the genre DRAMA has 866.7K and COMEDY 666.5K titles



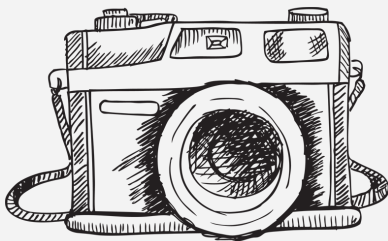


# separate features by modularized tables

Modular features which are on a per table basis:

- `df_ids = votes(df_ids)`
- `df_ids = principals(df_ids)`
- `df_ids = akas(df_ids)`
- `df_ids = basics(df_ids)`

**-----> We could have spent more time on the EDA, but we wanted to focus on the pipeline and wanted to sharpen our skills there.**



# Model Selection and Tuning

# XGB performs best

IMDb

2

01

## Linear Learner

Started with Linear Regression  
as a baseline model

1.32

objective\_type = 'Minimize'  
strategy='Random'

1.32

02

## KNN

KNN as the idea of nearest neighbor  
came across

as KNN brought the worst score, no  
further steps

1.44

03

## XGBoost

Trying to push weak learner features

1.22

objective\_type='Minimize'  
strategy='Bayesian'

1.20

Model

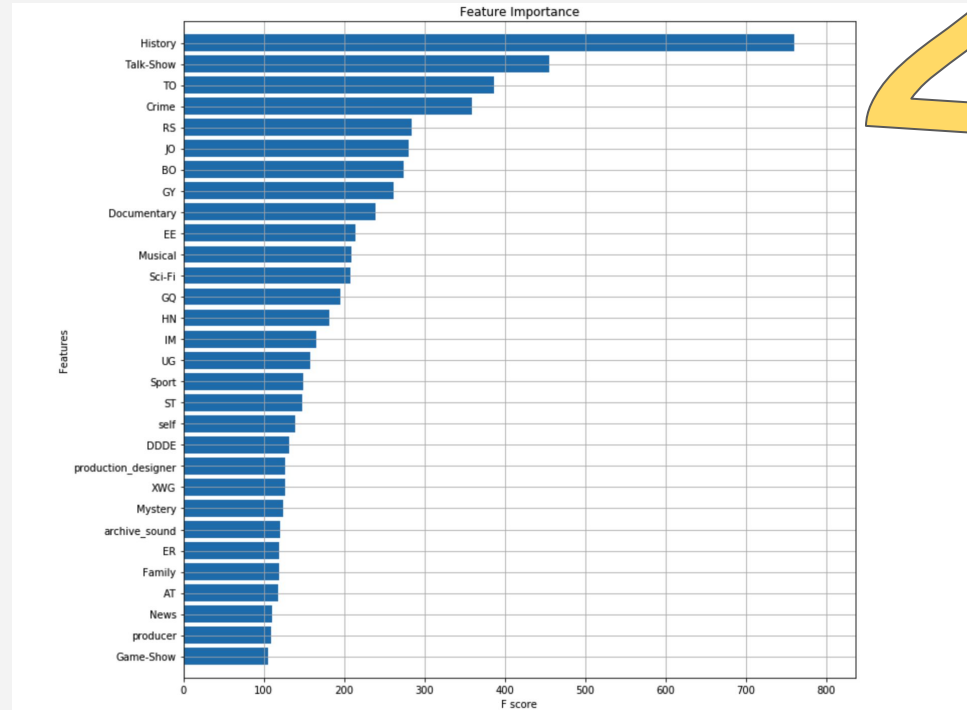
Test  
RMSE

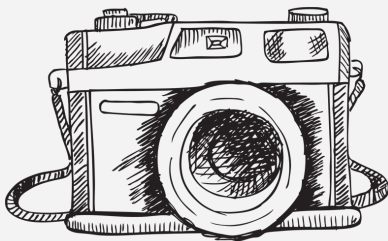
Hyper  
Parameter  
Tuner

# Model impacted by genre and country

First thoughts:

- History with biggest impact on features and it seems like that history movies have better ratings and smaller standard deviation compared to the overall count
- Other genres like Talk-Show and Crime are also important. Further Analysis required as to why.
- Countries: It could be that countries with a low number of movies are easier to split the dataset with.

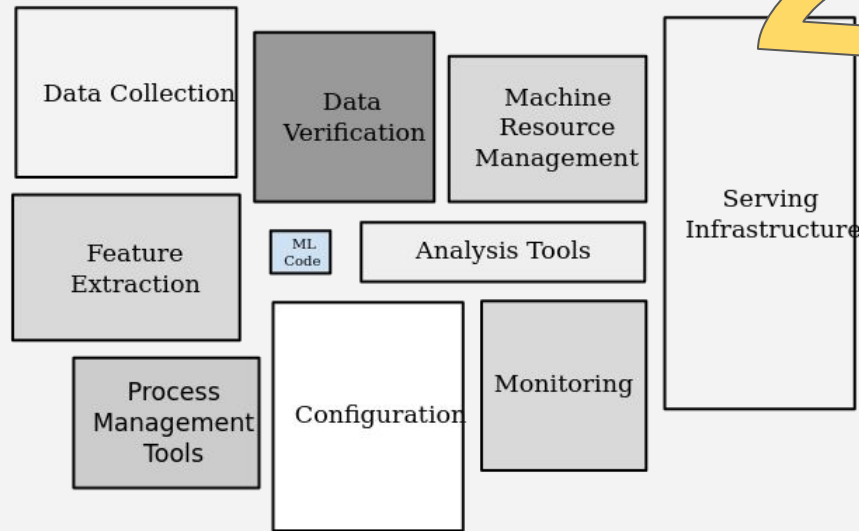




# Conclusion

# Lessons Learned

- Separate tasks (code cleaning, eda) can be conducted when waiting for a long query
- Always check the logic first (e.g. with a small sample) before conducting a query on the whole dataset
- Making different pieces of software talk to each other is complex



# Future Outlook

- Continue EDA
- Analyze Feature importance
- Automated Feature Selection
- Incorporate MLFlow
- Make the endpoint publicly available
- Implement logic which takes care of input data
- Document the model
- ...

[https://github.com/yaedin/imdb\\_project](https://github.com/yaedin/imdb_project)

Thank you!

